

Compulsory exercise 1: Group 4

TMA4268 Statistical Learning V2021

Øystein Alvestad, Lars Evje and William Scott Grundeland Olsen

16 februar, 2021

Problem 1

We consider the regression problem

$$Y = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p + \varepsilon,$$

where $E(\varepsilon) = 0$ and $\text{Var}(\varepsilon) = \sigma^2$. We define $\mathbf{x}, \boldsymbol{\beta} \in \mathbb{R}^{p+1}$ such that $f(\mathbf{x}) = \mathbf{x}^\top \boldsymbol{\beta} = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p$.

a)

We consider the estimator

$$\tilde{\boldsymbol{\beta}} = (X^\top X + \lambda I)^{-1} X^\top \mathbf{Y},$$

for $\boldsymbol{\beta}$. Here X is the design matrix, \mathbf{Y} is the response vector, I is the identity matrix, and $\lambda \geq 0$ is a constant. We recall that for a constant matrix A of appropriate dimensions, and a random vector \mathbf{Z} , we have that

$$E(A\mathbf{Z}) = A E(\mathbf{Z}) \quad \text{and} \quad \text{Var}(A\mathbf{Z}) = A E(\mathbf{Z}) A^\top. \quad (1)$$

First we now find the expected value

$$E(\tilde{\boldsymbol{\beta}}) = E((X^\top X + \lambda I)^{-1} X^\top \mathbf{Y}) = (X^\top X + \lambda I)^{-1} X^\top E(\mathbf{Y}),$$

and because $E(\mathbf{Y}) = X\boldsymbol{\beta}$, we get that

$$E(\tilde{\boldsymbol{\beta}}) = (X^\top X + \lambda I)^{-1} X^\top X \boldsymbol{\beta}.$$

Similarly, the variance-covariance matrix is

$$\text{Var}(\tilde{\boldsymbol{\beta}}) = \text{Var}((X^\top X + \lambda I)^{-1} X^\top \mathbf{Y}) = (X^\top X + \lambda I)^{-1} X^\top \text{Var}(\mathbf{Y}) [(X^\top X + \lambda I)^{-1} X^\top]^\top,$$

and because $\text{Var}(\mathbf{Y}) = \sigma^2 I$, we get that

$$\text{Var}(\tilde{\boldsymbol{\beta}}) = \sigma^2 (X^\top X + \lambda I)^{-1} X^\top X (X^\top X + \lambda I)^{-1},$$

where we used $(B^{-1})^\top = (B^\top)^{-1}$, for an invertible matrix B .

b)

We now let $\tilde{f}(\mathbf{x}_0) = \mathbf{x}_0^\top \tilde{\boldsymbol{\beta}}$ be the prediction at a new covariate vector \mathbf{x}_0 , and we wish the expected value and variance of this. Using what we learned in **a)** and Equation (1), we find that

$$E(\tilde{f}(\mathbf{x}_0)) = E(\mathbf{x}_0^\top \tilde{\boldsymbol{\beta}}) = \mathbf{x}_0^\top E(\tilde{\boldsymbol{\beta}}) = \mathbf{x}_0^\top (X^\top X + \lambda I)^{-1} X^\top X \boldsymbol{\beta}.$$

Similarly, we find that

$$\text{Var}(\tilde{f}(\mathbf{x}_0)) = \text{Var}(\mathbf{x}_0^\top \tilde{\boldsymbol{\beta}}) = \mathbf{x}_0^\top \text{Var}(\tilde{\boldsymbol{\beta}}) \mathbf{x}_0 = \sigma^2 \mathbf{x}_0^\top (X^\top X + \lambda I)^{-1} X^\top X (X^\top X + \lambda I)^{-1} \mathbf{x}_0. \quad (2)$$

c)

We want to find the expected MSE at \mathbf{x}_0 , and this is most easily done using the relationship between the expected value and the variance-covariance matrix,

$$E((y_0 - \tilde{f}(\mathbf{x}_0))^2) = E(\tilde{f}(\mathbf{x}_0) - f(\mathbf{x}_0))^2 + \text{Var}(\tilde{f}(\mathbf{x}_0)) + \text{Var}(\varepsilon).$$

The last two terms are known to us now, so we look further at the first term

$$\begin{aligned} E(\tilde{f}(\mathbf{x}_0) - f(\mathbf{x}_0))^2 &= [E(\tilde{f}(\mathbf{x}_0)) - E(f(\mathbf{x}_0))]^2 = [E(\tilde{f}(\mathbf{x}_0)) - f(\mathbf{x}_0)]^2 \\ &= [\mathbf{x}_0^\top (X^\top X + \lambda I)^{-1} X^\top X \beta - \mathbf{x}_0^\top \beta]^2. \end{aligned} \quad (3)$$

We then get that

$$\begin{aligned} E((y_0 - \tilde{f}(\mathbf{x}_0))^2) &= \sigma^2 + [\mathbf{x}_0^\top (X^\top X + \lambda I)^{-1} X^\top X \beta - \mathbf{x}_0^\top \beta]^2 \\ &\quad + \sigma^2 \mathbf{x}_0^\top (X^\top X + \lambda I)^{-1} X^\top X (X^\top X + \lambda I)^{-1} \mathbf{x}_0. \end{aligned}$$

d)

We start by importing the relevant quantities, as given in the project description.

```
id <- "1X_80KcoYbnglXvYFDirxjEW7LtpNr1m" # Google file ID
values <- dget(sprintf("https://docs.google.com/uc?id=%s&export=download", id))

X <- values$X
x0 <- values$x0
beta <- values$beta
sigma <- values$sigma
```

We may then calculate the squared bias $E(\tilde{f}(\mathbf{x}_0) - f(\mathbf{x}_0))^2$, using Equation (3). Plotting this for $\lambda \in [0, 2]$, we get the result as in Figure 1. We see that the squared bias increases with large $\lambda > 0.5$, and has a minimum when $\lambda = 0$ and $\lambda \approx 0.5$. This is expected for $\lambda = 0$, because we then get that $\tilde{\beta}$ is equal to the OLS estimator. The bias measures how good the estimator is in estimating the real value, so this makes sense. From the figure it also appears that for $\lambda \approx 0.5$, the estimator $\tilde{\beta}$ is estimating the real value β very good. For increasing λ after this, the estimator seems to do a worse job in estimating β .

```
library(ggplot2)
bias <- function(lambda, X, x0, beta) {
  p <- ncol(X)
  inv <- solve(t(X) %*% X + lambda * diag(p))
  value <- (t(x0) %*% inv %*% t(X) %*% X %*% beta - t(x0) %*% beta)^2
  return(value)
}

lambdas <- seq(0, 2, length.out = 500)
BIAS <- rep(NA, length(lambdas))

for (i in 1:length(lambdas)) {
  BIAS[i] <- bias(lambdas[i], X, x0, beta)
}

dfBias <- data.frame(lambdas = lambdas, bias = BIAS)

ggplot(dfBias, aes(x = lambdas, y = bias)) +
  geom_line(color = "red") +
  xlab(expression(lambda)) +
  ylab(expression(bias^2))
```




Figure 1: The squared bias $E(\tilde{f}(\mathbf{x}_0) - f(\mathbf{x}_0))^2$ as a function of λ .

e)

We are now interested in $\text{Var}(\tilde{f}(\mathbf{x}_0))$, which can be calculated using Equation (2). Using $\lambda \in [0, 2]$, the result is plotted in Figure 2. We notice that the variance is decreasing for increasing values of λ . That is, as long as λ increases, the model is less prone to overfitting. Letting $\lambda \rightarrow \infty$, we also see that the variance tends to zero.

```
variance <- function(lambda, X, x0, sigma) {
  p <- ncol(X)
  inv <- solve(t(X) %*% X + lambda * diag(p))
  value <- sigma^2 * t(x0) %*% inv %*% t(X) %*% X %*% inv %*% x0
  return(value)
}

lambdas <- seq(0, 2, length.out = 500)
VAR <- rep(NA, length(lambdas))

for (i in 1:length(lambdas)) {
  VAR[i] <- variance(lambdas[i], X, x0, sigma)
}
dfVar <- data.frame(lambdas = lambdas, var = VAR)

ggplot(dfVar, aes(x = lambdas, y = var)) +
  geom_line(color = "green4") +
  xlab(expression(lambda)) +
  ylab("variance") +
  theme(plot.title = element_text(hjust = 0.5))
```

f)

Lastly, we are interested in the expected MSE at \mathbf{x}_0 , $E((y_0 - \tilde{f}(\mathbf{x}_0))^2)$, which, when we know the squared bias and the variance from e), is given as $(\text{bias})^2 + \text{Var}(\tilde{f}(\mathbf{x}_0)) + \sigma^2$, because the irreducible error is σ^2 . The plot of the expected MSE, the squared bias and the variance is shown in Figure 3, and by using `lambdas[which.min(exp_mse)]` we find that the minimal expected MSE is found when $\lambda \approx 0.993988$. That




Figure 2: The variance $\text{Var}(\tilde{f}(\mathbf{x}_0))$ as a function of λ .

is, it is possible to move to $\lambda > 0$, and taking a little bias, and reducing the variance, leading to a reduction in the expected MSE.

```
exp_mse <- BIAS + VAR + sigma^2

cols <- c("exp_mse" = "blue", "bias" = "red", "variance" = "green4")
dfAll <- data.frame(lambda = lambdas, bias = BIAS, var = VAR, exp_mse = exp_mse)

ggplot(dfAll)+
  geom_line(aes(x = lambda, y = exp_mse, color = "exp_mse")) +
  geom_line(aes(x = lambda, y = bias, color = "bias")) +
  geom_line(aes(x = lambda, y = var, color = "variance")) +
  xlab(expression(lambda)) +
  ylab(expression(E(MSE))) +
  theme(legend.title = element_blank()) +
  scale_color_manual(values = cols)
```




Figure 3: The expected MSE $E((y_0 - \tilde{f}(\mathbf{x}_0))^2)$ as a function of λ , together with the squared bias and the variance.

Problem 2

- a)
- b)
- c)
- d)

Problem 3

- a)
- b)
- c)
- d)

Problem 4

- a)
- b)

Problem 5

- a)
- b)