

Compulsory exercise 2: Group 4

TMA4268 Statistical Learning V2021

Øystein Alvestad, Lars Evje and William Scott Grundeland Olsen

16 april, 2021

Problem 1

- a)
- b)
- c)
- d)
- e)
- f)

Problem 2

- a)
- b)
- c)
- (i)
- (ii)

Problem 3

- a)

True, True, False, False

b)

c)

(i)

```
id <- "1Fv6xwKLSZHldRAC1MrcK2mzd0Ynbgv0E" # google file ID
d.diabetes <- dget(sprintf("https://docs.google.com/uc?id=%s&export=download", id))

d.train = d.diabetes$ctrain
d.train$diabetes = as.factor(d.train$diabetes)

d.test = d.diabetes$ctest
d.test$diabetes = as.factor(d.test$diabetes)
```

```
set.seed(1)
t.diabetes = tree(diabetes ~ ., data = d.train)

t.diabetes.pred = predict(t.diabetes, d.test, type = "class")
misclass = table(t.diabetes.pred, d.test$diabetes)
misclass
```

```
##
## t.diabetes.pred    0    1
##                0 126  28
##                1  29  49
```


```
1 - sum(diag(misclass)/sum(misclass))
```

```
## [1] 0.2456897
```

```
cv.diabetes = cv.tree(t.diabetes, FUN = prune.misclass, K = 10)
cv.diabetes
```

```
## $size
## [1] 17 14 10  6  2  1
##
## $dev
## [1]  77  78  85  81  82 100
##
## $k
## [1] -Inf  0.00  1.50  2.75  5.00 29.00
##
## $method
## [1] "misclass"
##
## attr(,"class")
## [1] "prune"          "tree.sequence"
```

```
plot(cv.diabetes$size, cv.diabetes$dev, type = "b", xlab = "Terminal nodes")
```



```
prune.diabetes = prune.misclass(t.diabetes, best = 3)
# plot(prune.diabetes) text(prune.diabetes, pretty = 1)

prune.diabetes.pred = predict(prune.diabetes, d.test, type = "class")
misclass.prune = table(prune.diabetes.pred, d.test$diabetes)
misclass.prune
```

```
##
## prune.diabetes.pred  0  1
##                   0 119 20
##                   1  36 57
```

```
1 - sum(diag(misclass.prune)/sum(misclass.prune))
```

```
## [1] 0.2413793
```

The 10-fold CV cost-complexity, with `FUN = prune.misclass` tells us to choose the full tree. Using `FUN = deviance` however gives us that we should choose the tree with 3 terminal nodes. The misclassification error for the full tree is 0.2457, while for the tree with three terminal nodes it is 0.2414 so slightly better on the test set.

(ii)

```
library(randomForest)
set.seed(1)
rf.diabetes = randomForest(diabetes ~ ., data = d.train, ntree = 1000, mtry = 3,
                           importance = TRUE)

rf.diabetes.pred = predict(rf.diabetes, d.test, type = "class")
misclass.rf = table(rf.diabetes.pred, d.test$diabetes)
misclass.rf
```

```
##
## rf.diabetes.pred    0    1
##                   0 135  34
##                   1  20  43
```

```
1 - sum(diag(misclass.rf)/sum(misclass.rf))
```

```
## [1] 0.2327586
```

```
rf.diabetes
```

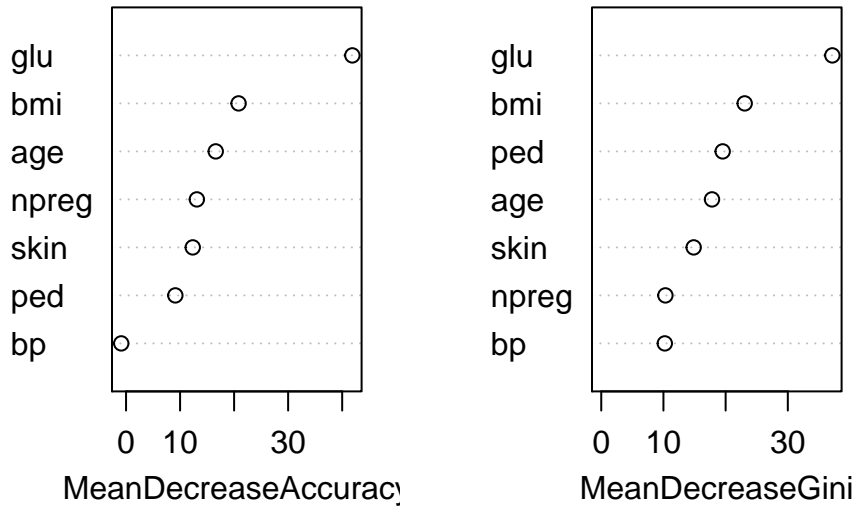
```
##
## Call:
## randomForest(formula = diabetes ~ ., data = d.train, ntree = 1000, mtry = 3, importance = TRUE,
##               Type of random forest: classification
##               Number of trees: 1000
## No. of variables tried at each split: 3
##
## OOB estimate of error rate: 20%
## Confusion matrix:
##    0  1 class.error
## 0 175 25      0.125
## 1  35 65      0.350
```

```
importance(rf.diabetes)
```

```
##              0              1 MeanDecreaseAccuracy MeanDecreaseGini
## npreg 14.968249  0.9760653      13.1023756      10.31516
## glu   33.655803 29.8766042      41.8842028      37.16251
## bp     1.995766 -3.5518898      -0.8901241      10.22150
## skin  10.536315  5.6018397      12.3438407      14.86193
## bmi   15.489014 14.4753432      20.8293113      23.06560
## ped    7.620861  4.8751738       9.1223090      19.52611
## age   15.402886  5.9624002      16.5960946      17.81053
```

```
varImpPlot(rf.diabetes)
```

rf.diabetes



We use the more advanced method Random Forest, with the tuning parameters `ntree=1000`, and `mtry=3`. From the importance plots we observe that `glu` and `bmi` is the most influential variables in the prediction of diabetes.

Problem 4

a)

False, True, False, True.

b)

(i)

A SVM tends to behave better than logistic regression when the classes are well separated, which we would assume them to be here with the genomic data. Also SVM has the ability to work in high dimensional space compared to the small number of samples.

Instead of SVM one could for example use a random forest, K -nearest neighbor, linear classifiers, or quadratic classifiers.

(ii)

The paper introduces an ensemble SVM-Recursive Feature Elimination for gene selection that follows the concept of ensemble and bagging used in random forest but adopts the backward elimination strategy which is the rationale of Recursive Feature Elimination algorithm.

(iii)

In the following code block we fit a support vector classifier with $C = 1$ on `Category` using `d.leukemia.train`.

```
set.seed(2399)

# Set-up:
id <- "1x_E8xnmz9CMHh_tMwIsWP94czPa1Fpsj" # Google file ID
path <- "https://docs.google.com/uc?id=%s&export=download"
d.leukemia <- read.csv(sprintf(path, id), header = TRUE)

t.samples <- sample(1:60, 15, replace = FALSE)
d.leukemia$Category <- as.factor(d.leukemia$Category)
d.leukemia.test <- d.leukemia[t.samples, ]
d.leukemia.train <- d.leukemia[-t.samples, ]

# Support vector classifier:
svmfit <- svm(Category ~ ., data = d.leukemia.train, kernel = "linear", cost = 1,
  scale = TRUE)
pred_train <- predict(svmfit, d.leukemia.train)
pred_test <- predict(svmfit, d.leukemia.test)

# Confusion tables for training and testing respectively:
conf_tab_train <- table(predict = pred_train, truth = d.leukemia.train$Category)
conf_tab_test <- table(predict = pred_test, truth = d.leukemia.test$Category)

# Misclassification for training and testing respectively:
misclas_train <- 1 - sum(diag(conf_tab_train))/sum(conf_tab_train)
misclas_test <- 1 - sum(diag(conf_tab_test))/sum(conf_tab_test)
```

We then see the confusion table for the training data below, with the misclassification error rate of 0.

```
conf_tab_train
```

```
##           truth
## predict   Non-Relapse Relapse
## Non-Relapse      30      0
## Relapse          0      15
```

We also have the confusion table for the test data below, with the misclassification error rate of 0.3333333.

```
conf_tab_test
```

```
##           truth
## predict   Non-Relapse Relapse
## Non-Relapse      8      4
## Relapse          1      2
```

The training error rate is 0, suggesting that there is an overfitting of the data. This is dependent on the cost C , and one could have done a cross validation to find a possibly better cost than $C = 1$.

The most common error in the test set is that the truth is relapse, while the prediction is non-relapse. That is, children relapse even though the prediction is that they do not. With a misclassification error rate of 0.333333 for the test set the classification can be said to be successful. However, the false positive, which is the most common error, is worse than the false negative in this case, in our opinion.

(iv)

In the following code block we fit a support vector machine to the data using the cost $C = 1$ and the tuning parameter $\gamma = 10^{-2}$ or $\gamma = 10^{-5}$.

```
set.seed(2399)

# Support vector machine and prediction:
svmfit_gamma1 <- svm(Category ~ ., data = d.leukemia.train, kernel = "radial", cost = 1,
  gamma = 0.01, scale = TRUE)
svmfit_gamma2 <- svm(Category ~ ., data = d.leukemia.train, kernel = "radial", cost = 1,
  gamma = 1e-05, scale = TRUE)
pred_train_gamma1 <- predict(svmfit_gamma1, d.leukemia.train)
pred_test_gamma1 <- predict(svmfit_gamma1, d.leukemia.test)
pred_train_gamma2 <- predict(svmfit_gamma2, d.leukemia.train)
pred_test_gamma2 <- predict(svmfit_gamma2, d.leukemia.test)

# Confusion tables for training and testing:
conf_tab_train_gamma1 <- table(predict = pred_train_gamma1, truth = d.leukemia.train$Category)
conf_tab_test_gamma1 <- table(predict = pred_test_gamma1, truth = d.leukemia.test$Category)
conf_tab_train_gamma2 <- table(predict = pred_train_gamma2, truth = d.leukemia.train$Category)
conf_tab_test_gamma2 <- table(predict = pred_test_gamma2, truth = d.leukemia.test$Category)

# Misclassification for training and testing:
misclas_train_gamma1 <- 1 - sum(diag(conf_tab_train_gamma1))/sum(conf_tab_train_gamma1)
misclas_test_gamma1 <- 1 - sum(diag(conf_tab_test_gamma1))/sum(conf_tab_test_gamma1)
misclas_train_gamma2 <- 1 - sum(diag(conf_tab_train_gamma2))/sum(conf_tab_train_gamma2)
misclas_test_gamma2 <- 1 - sum(diag(conf_tab_test_gamma2))/sum(conf_tab_test_gamma2)
```

We then see the confusion table for the training data for $\gamma = 10^{-2}$ below, with the misclassification error rate of 0.

```
conf_tab_train_gamma1
```

```
##           truth
## predict    Non-Relapse Relapse
## Non-Relapse      30      0
## Relapse          0     15
```

We also have the confusion table for the test data for $\gamma = 10^{-2}$ below, with the misclassification error rate of 0.4.

```
conf_tab_test_gamma1
```

```
##           truth
## predict    Non-Relapse Relapse
## Non-Relapse      9      6
## Relapse          0      0
```

For $\gamma = 10^{-5}$ we have the confusion table for the training data below, with the misclassification error rate of 0.333333.

```
conf_tab_train_gamma2
```

```
##           truth
## predict   Non-Relapse Relapse
## Non-Relapse      30      15
## Relapse          0       0
```

We also have the confusion table for the test data for $\gamma = 10^{-5}$ below, with the misclassification error rate of 0.4.

```
conf_tab_test_gamma2
```

```
##           truth
## predict   Non-Relapse Relapse
## Non-Relapse      9       6
## Relapse          0       0
```

We note that the misclassification error rate for the training set is 0 for $\gamma = 10^{-2}$ and 0.333333 for $\gamma = 10^{-5}$. This can be explained by the fact that for small γ the decision boundaries are smoother than for larger γ . Thus, there may be some overfitting for $\gamma = 10^{-2}$. For the test data however, the results are the same. Comparing to the case in (iii), the results are worse, suggesting that the support vector classifier is better than the support vector machine for this dataset.

c)

The polynomial kernel of positive integer degree d has the form

$$K(\mathbf{x}, \mathbf{y}) = (1 + \mathbf{x}^\top \mathbf{y})^d = \left(1 + \sum_{i=1}^p x_i y_i\right)^d,$$

for $\mathbf{x}, \mathbf{y} \in \mathbb{R}^p$, with elements x_i and y_i for $i = 1, \dots, p$. We assume $d = 2$ and $\mathbf{x}, \mathbf{y} \in \mathbb{R}^2$, such that

$$\begin{aligned} K(\mathbf{x}, \mathbf{y}) &= (1 + \mathbf{x}^\top \mathbf{y})^2 = 1 + 2\mathbf{x}^\top \mathbf{y} + (\mathbf{x}^\top \mathbf{y})^2 = 1 + 2(x_1 y_1 + x_2 y_2) + (x_1 y_1 + x_2 y_2)^2 \\ &= 1 + 2x_1 y_1 + 2x_2 y_2 + x_1^2 y_1^2 + x_2^2 y_2^2 + 2x_1 y_1 x_2 y_2. \end{aligned}$$

We then see that

$$K(\mathbf{x}, \mathbf{y}) = \mathbf{h}(\mathbf{x})^\top \mathbf{h}(\mathbf{y}) = \langle \mathbf{h}(\mathbf{x}), \mathbf{h}(\mathbf{y}) \rangle,$$

by the basic definition of the inner product of two vectors, where,

$$\mathbf{h}(\mathbf{x}) = \begin{bmatrix} 1 \\ \sqrt{2}x_1 \\ \sqrt{2}x_2 \\ x_1^2 \\ x_2^2 \\ \sqrt{2}x_1 x_2 \end{bmatrix} \quad \text{and} \quad \mathbf{h}(\mathbf{y}) = \begin{bmatrix} 1 \\ \sqrt{2}y_1 \\ \sqrt{2}y_2 \\ y_1^2 \\ y_2^2 \\ \sqrt{2}y_1 y_2 \end{bmatrix}.$$

Problem 5

a)

True, False, False, False.

b)

In the following we make a random cluster of the data and compute the centroid of the two clusters we get. The clusters are color coded where one cluster is colored red, while the other is green, as seen in Figure 1. Note that the code here is not general for every K -mean clustering, but is only applicable to $K = 2$, which is the case given in the problem.

```
set.seed(1)

x1 <- c(1, 2, 0, 4, 5, 6)
x2 <- c(5, 4, 3, 1, 1, 2)
X <- matrix(c(x1, x2), ncol = 2)

# Random cluster
X_cluster <- cbind(X, sample(c(1, 2), size = nrow(X), replace = TRUE))

# Initializing and computing the centroids:
g1_centroid <- c(0, 0)
g2_centroid <- c(0, 0)

for (i in 1:length(x1)) {
  if (X_cluster[i, 3] == 1) {
    g1_centroid[1] <- g1_centroid[1] + X[i, 1]
    g1_centroid[2] <- g1_centroid[2] + X[i, 2]
  } else {
    g2_centroid[1] <- g2_centroid[1] + X[i, 1]
    g2_centroid[2] <- g2_centroid[2] + X[i, 2]
  }
}

g1_centroid <- g1_centroid/length((X[, 1])[X_cluster[, 3] == 1])
g2_centroid <- g2_centroid/length((X[, 1])[X_cluster[, 3] == 2])

# Plotting the clusters and centroids color coded:
plot(X, col = X_cluster[, 3] + 1, main = "Random clustering of the data with the centroids",
     xlab = "x1", ylab = "x2", pch = 20, cex = 2)
points(g1_centroid[1], g1_centroid[2], pch = 15, cex = 2, col = 2)
points(g2_centroid[1], g2_centroid[2], pch = 15, cex = 2, col = 3)
```

We can then measure, using the Euclidean distance, what points are closest to the respective centroids, in this case giving the correct clustering for $K = 2$. This is shown in Figure 2.

```
dist <- function(x, y) {
  return(sqrt(sum((x - y)^2)))
}
```

Random clustering of the data with the centroids




Figure 1: A random clustering of the data being the round points, and the centroids being the square points.

```
for (i in 1:length(x1)) {  
  X_cluster[i, 3] <- ifelse(dist(g1_centroid, X[i, ]) < dist(g2_centroid, X[i,  
    ]), 1, 2)  
}  
  
plot(X, col = X_cluster[, 3] + 1, main = "K-means clustering of the data with K = 2",  
  xlab = "x1", ylab = "x2", pch = 20, cex = 2)
```

```
id <- "1VfVCQvWt121UN39NXZ4aR9Dmsbj-p90U" # google file ID  
GeneData <- read.csv(sprintf("https://docs.google.com/uc?id=%s&export=download",  
  id), header = F)  
colnames(GeneData)[1:20] = paste(rep("H", 20), c(1:20), sep = "")  
colnames(GeneData)[21:40] = paste(rep("D", 20), c(1:20), sep = "")  
row.names(GeneData) = paste(rep("G", 1000), c(1:1000), sep = "")  
GeneData = t(GeneData)  
GeneData <- scale(GeneData)
```

K-means clustering of the data with $K = 2$


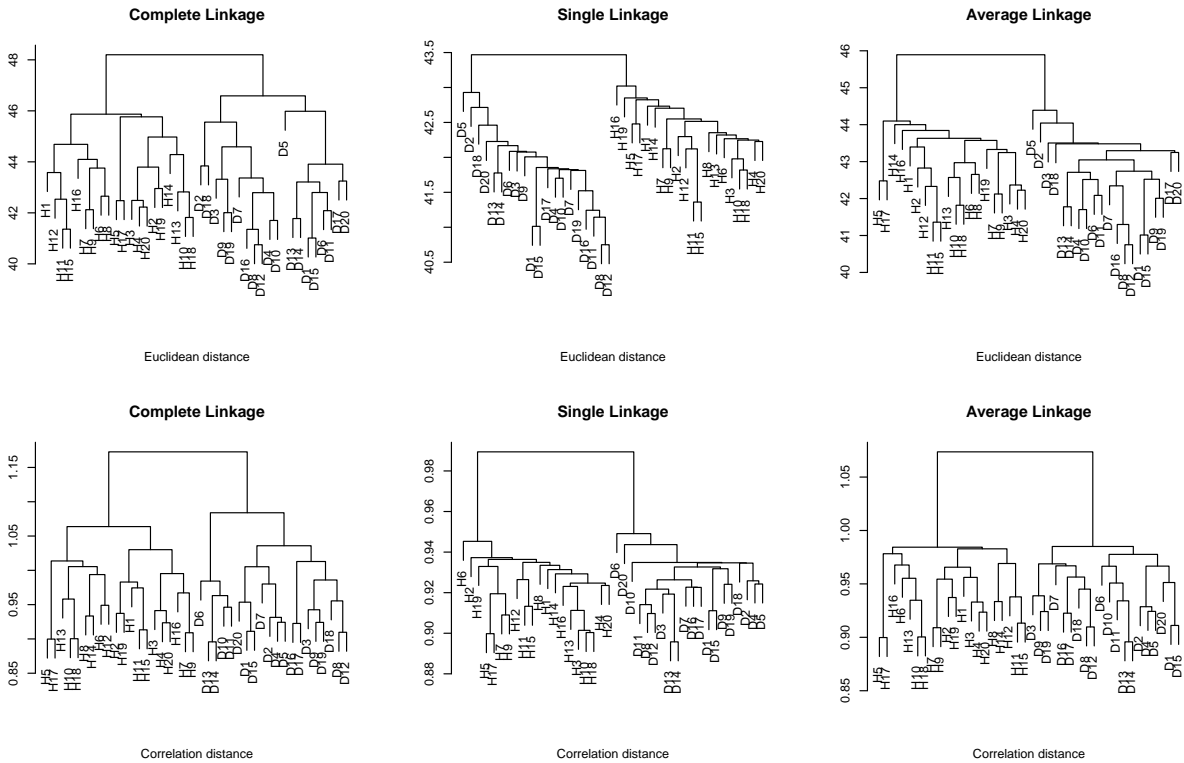


Figure 2: The K -means clustering for the data, with $K = 2$.

c)



d)

```
cutree(hc.Gene.EcC, 2)
```

```
## H1 H2 H3 H4 H5 H6 H7 H8 H9 H10 H11 H12 H13 H14 H15 H16 H17 H18 H19 H20
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## D1 D2 D3 D4 D5 D6 D7 D8 D9 D10 D11 D12 D13 D14 D15 D16 D17 D18 D19 D20
## 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
cutree(hc.Gene.EcS, 2)
```

```
## H1 H2 H3 H4 H5 H6 H7 H8 H9 H10 H11 H12 H13 H14 H15 H16 H17 H18 H19 H20
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## D1 D2 D3 D4 D5 D6 D7 D8 D9 D10 D11 D12 D13 D14 D15 D16 D17 D18 D19 D20
## 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
cutree(hc.Gene.EcA, 2)
```

```
## H1 H2 H3 H4 H5 H6 H7 H8 H9 H10 H11 H12 H13 H14 H15 H16 H17 H18 H19 H20
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## D1 D2 D3 D4 D5 D6 D7 D8 D9 D10 D11 D12 D13 D14 D15 D16 D17 D18 D19 D20
## 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
cutree(hc.Gene.CorrC, 2)
```

```
## H1 H2 H3 H4 H5 H6 H7 H8 H9 H10 H11 H12 H13 H14 H15 H16 H17 H18 H19 H20
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## D1 D2 D3 D4 D5 D6 D7 D8 D9 D10 D11 D12 D13 D14 D15 D16 D17 D18 D19 D20
## 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
cutree(hc.Gene.CorrS, 2)
```

```
## H1 H2 H3 H4 H5 H6 H7 H8 H9 H10 H11 H12 H13 H14 H15 H16 H17 H18 H19 H20
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## D1 D2 D3 D4 D5 D6 D7 D8 D9 D10 D11 D12 D13 D14 D15 D16 D17 D18 D19 D20
## 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
cutree(hc.Gene.CorrA, 2)
```

```
## H1 H2 H3 H4 H5 H6 H7 H8 H9 H10 H11 H12 H13 H14 H15 H16 H17 H18 H19 H20
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## D1 D2 D3 D4 D5 D6 D7 D8 D9 D10 D11 D12 D13 D14 D15 D16 D17 D18 D19 D20
## 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Since it was given that the 20 first tissue samples was healthy, and the last 20 was damaged we observe that all of the above hierarchical clusterings managed to separate the two groups perfectly.

e)


(i)

```
pc.Gene = prcomp(GeneData)
summary(pc.Gene)
```


```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation    9.00460 5.87302 5.74347 5.61806 5.55344 5.50107 5.40069
## Proportion of Variance 0.08108 0.03449 0.03299 0.03156 0.03084 0.03026 0.02917
## Cumulative Proportion 0.08108 0.11558 0.14856 0.18013 0.21097 0.24123 0.27040
##              PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation    5.38575 5.3762 5.34146 5.31878 5.25016 5.18737 5.1667
## Proportion of Variance 0.02901 0.0289 0.02853 0.02829 0.02756 0.02691 0.0267
## Cumulative Proportion 0.29940 0.3283 0.35684 0.38513 0.41269 0.43960 0.4663
##              PC15     PC16     PC17     PC18     PC19     PC20     PC21
## Standard deviation    5.10384 5.04667 5.03288 4.98926 4.92635 4.90996 4.88803
## Proportion of Variance 0.02605 0.02547 0.02533 0.02489 0.02427 0.02411 0.02389
## Cumulative Proportion 0.49234 0.51781 0.54314 0.56803 0.59230 0.61641 0.64030
##              PC22     PC23     PC24     PC25     PC26     PC27     PC28
## Standard deviation    4.85159 4.79974 4.78202 4.70171 4.66105 4.64595 4.59194
## Proportion of Variance 0.02354 0.02304 0.02287 0.02211 0.02173 0.02158 0.02109
## Cumulative Proportion 0.66384 0.68688 0.70975 0.73185 0.75358 0.77516 0.79625
##              PC29     PC30     PC31     PC32     PC33     PC34     PC35
## Standard deviation    4.53246 4.47381 4.4389 4.41670 4.39404 4.3591 4.23504
## Proportion of Variance 0.02054 0.02001 0.0197 0.01951 0.01931 0.0190 0.01794
## Cumulative Proportion 0.81679 0.83681 0.8565 0.87602 0.89533 0.9143 0.93226
##              PC36     PC37     PC38     PC39     PC40
## Standard deviation    4.2184 4.12936 4.0738 4.03658 5.5e-15
## Proportion of Variance 0.0178 0.01705 0.0166 0.01629 0.0e+00
## Cumulative Proportion 0.9501 0.96711 0.9837 1.00000 1.0e+00
```

```
cols = c(rep("green", 20), rep("red", 20))
```

```
plot(pc.Gene$x[, 1:2], col = cols)
```



```
GeneLoad = pc.Gene$rotation[, 1:2]
informative_loadings = rbind(GeneLoad[order(GeneLoad[, 1], decreasing = TRUE), ][1:5, ],
                             GeneLoad[order(GeneLoad[, 2], decreasing = TRUE), ][1:5, ])
biplot(x = pc.Gene$x[, 1:2], y = informative_loadings, scale = 0)
```



```
tail(sort(abs(pc.Gene$rotation[, 1])), 20)
```

```
##      G508      G540      G564      G566      G592      G528      G535
## 0.08541371 0.08550741 0.08552071 0.08553015 0.08558608 0.08561425 0.08610094
##      G599      G570      G511      G509      G584      G538      G593
## 0.08624312 0.08626458 0.08655126 0.08661015 0.08690858 0.08745400 0.08758616
##      G551      G600      G590      G565      G589      G502
## 0.08768360 0.09167322 0.09173169 0.09183823 0.09449766 0.09485044
```

(ii)

From the summary above we read of that the first 5 principle components explain around 21.1% of the variance.

f)

From the plot above we see that the two groups can be separated by looking at the first principle component alone in this case, in addition by the properties of PCA we know that PC1 capture most of the variance (out of any PC). Therefore we look at the genes with the highest loadings in PC1, which gives us that